

AI Pilgrimage · Rail Rebirth Belt — Jing-Zhang Auditable Scenario Line

- English name: Jing-Zhang AI Pilgrimage & Proof Line
- Core system: JZ-AIOS (Jing-Zhang Auditable Innovation Operating System)

Design Basis and Source List

This second-round plan begins by acknowledging a fact: although the existing proposal passed the 8/8 baseline score and the spatial, visual, and professional self-checks, “validation passed” only means that the files are complete, the format is correct, and the references exist. It does not prove that the proposal has established a distinctive urban method. V1’s “one spine, three anchors, six interfaces, and multiple nodes” remained close to a conventional corridor plan: the three anchors were positioned but the two wings were absent; the ten AI scenarios were described in only one sentence each; three public-space nodes were insufficient to support continuous public life; and the task matrix repeatedly reused the same evidence. V2 therefore stops adding isolated applications and instead consolidates space, industry, culture, scenarios, and operations into “an auditable urban innovation production line” [depth:existing_conditions_diagnosis].

Current shortcoming	Verifiable evidence	V2 planning response
Three parallel anchors, with no operating loop across the three zones and two wings	The text positions three places but provides no relay for problem intake, verification, publication, and feedback	Establish a state machine: “Xiaoyue River problem intake → Origin Community co-creation → Zhongzhi verification → Dazhongsi publication → Zhongguancun transformation → public feedback”
Ten scenarios are a feature list rather than operating protocols	scenario_count is derived only by counting numbered items in the text	Establish 12 machine-readable scenario nodes; the ten service scenarios plus two public-audit nodes all receive passports, risk controls, manual fallbacks, expiry, and rollback
The six spatial interfaces are slogans only	No Feature IDs, locations, functions, or acceptance criteria	Materialize the six spatial interfaces as PUBLIC-004—PUBLIC-009 and bind them to lateral stitching routes, non-AI access, quiet hours, and the right to pause
Spatial expression in the key areas is too sparse	V1 contains only a few randomly placed masses and cannot show block–frontage–scenario relationships	Concentrate and refine the massing prototypes in the three provisional key areas, adding three attributes: retain for assessment, adaptive-reuse candidate, and new-build candidate
Public-space and land-use measurement scopes conflict	V1 has only three PUBLIC_SPACE nodes, while land-use code 1403 is a broad functional envelope	Distinguish “land-use functional envelope” from “direct design-intervention footprint” and expand to nine reversible public nodes
Metrics are overly geometric	No metrics for passports, manual fallback, non-AI access, or quiet and screen-free space	Expand the indicators into seven categories: space, governance, public value, industrial verification, cultural trustworthiness, operations, and resilience
Operations cover only organization and annual reporting	No annual rhythm, admission gates, outcome transformation, or public disclosure of failures	Establish a four-season protocol: Problem Season → Open-source Season → Urban Beta Season → Proof Week

The Real Starting Point after 6 August 2026

Public reporting states that Phase II of the Jing-Zhang Railway Heritage Park opened on 6 August 2026. Its southern community-vitality section and northern nature-and-leisure section now form a reported green corridor of about 9 km, with existing outcomes in community connectivity, all-age activity, slow mobility, and sponge infrastructure [source:HAIDIAN-JZ-PHASE2-OPEN-2026]. This changes the proposal’s basic judgment: JZ-AIOS no longer treats “building a continuous park” as its own outcome. It proposes only a reversible, low-footprint, and removable public-service and verification increment over the existing park. The table below is a design starting point derived from desk research, not a site visit, as-built drawing review, or confirmation of a statutory red line. Location, operating condition, and actual use still require item-by-item verification [data:visual/assets/site-grounding-register.json#SG-001] [assumption:A-SITE-LOCATION-008].

Publicly reported baseline	Design increment proposed here	To be verified on site
Phase II is open; the southern section is reported to serve adjacent communities and open lateral connections	Recast the six spatial interfaces as audit/service interface types that first reuse existing entrances, staffed services, and slow-mobility nodes rather than presuming new landmarks	Exact location, accessible continuity, peak-hour detours, ownership, and operations boundary at each interface
The northern section is reported to provide fishbone slow-mobility links and connections to waterfront and park networks	Add only pausable scenario passports, offline wayfinding, and public retest interfaces; do not rebuild the basic greenway	Conflicts among walking, cycling, and running; night lighting; stormwater and railway-safety conditions
All-age activity, heritage reuse, and sponge facilities have been publicly reported	AI must prove net value over everyday use; device count, screens, and event footfall cannot substitute for public value	A 30-day post-opening use audit, task success across groups, noise, complaints, and maintenance burden

Built-Park Delta Protocol (INNOV-006)

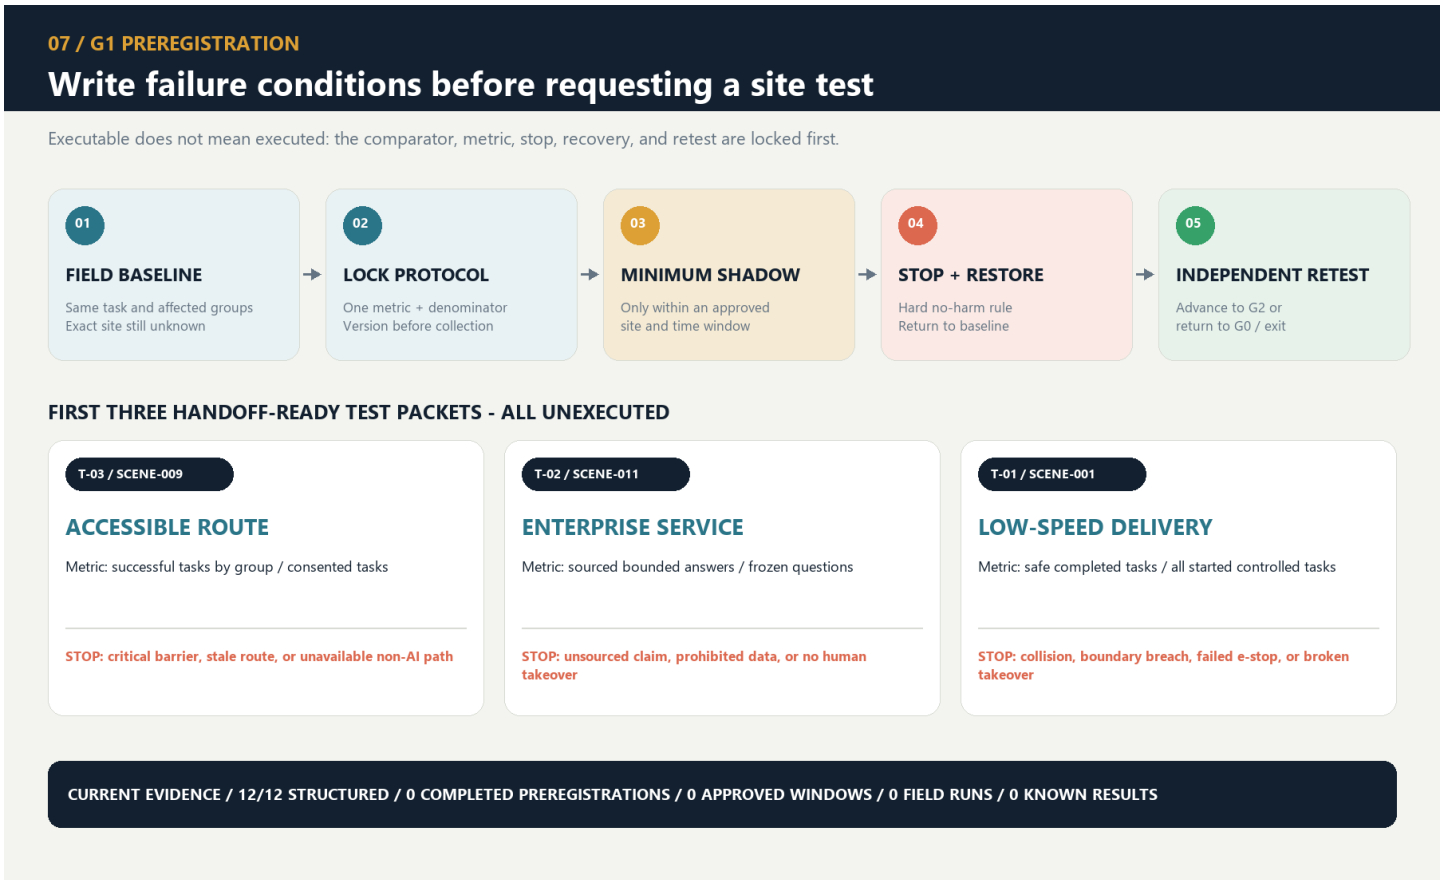
Before any G1/G2 pilot, one page must record the “non-AI current baseline → AI increment hypothesis → minimum footprint and time window → stop and recovery → independent retest” chain. AI must prove incremental value over staffed, paper-based, physical-wayfinding, or existing services, with no net loss of accessible public space or quiet periods. If it cannot prove an increment, cannot restore a state no worse than the baseline, or fails independent retesting, the pilot returns to G0 or exits. The completed 30-day post-opening field-audit count is currently 0; protocol completeness proves only design-field coverage, not field effectiveness [data:visual/assets/innovation-register.json#INNOV-006] [assumption:A-POST-OPENING-007] [metric:completed_post_opening_field_audit_count].

G1 Preregistration: Write the Failure Conditions before Requesting a Site Test

visual/assets/g1-preregistration-register.json turns all 12 existing scenarios into executable-but-unexecuted first-test contracts. Before testing, each record must lock a non-AI comparison task, one primary metric and denominator, affected groups, sample frame, collection window, no-harm rule, stop and recovery conditions, independent-retest inputs, and version. The 12/12 figure proves required-field coverage only. Completed preregistrations, approved G1 windows, field executions, and known results all remain 0, so no scene may advance from G0 on the strength of this register [data:visual/assets/g1-preregistration-register.json#G1-PREREG-001] [metric:g1_preregistration_required_field_coverage_ratio] [metric:completed_g1_preregistration_count].

Review path (a conceptual protocol, not approval or results): built-park baseline → public problem → same-task non-AI baseline → G1 shadow test → stop/recover → independent retest → upgrade or exit. Every node remains at G0; this path is not approval, field execution, or field results.

First minimum test packet	Same-task non-AI comparator	Single primary metric	Non-negotiable stop condition	Current status
T-03 / SCENE-009 accessible route	Staffed service, paper, or physical wayfinding for the same origin and destination	Successful tasks by co-test group / consented tasks	Stop and restore on a critical barrier, stale route state, or unavailable non-AI route	Threshold pending field preregistration; unexecuted
T-02 / SCENE-011 enterprise service	Staffed counter or traceable static guide using the same frozen question set	Sourced and correctly bounded answers / frozen questions	Stop on an unsourced conclusion, prohibited-data exposure, or unavailable human takeover	Threshold pending accountable and professional review; unexecuted
T-01 / SCENE-001 low-speed delivery	Staffed delivery or static route for the same task	Tasks with no collision, no boundary breach, and a working physical stop / approved controlled tasks	Stop on any collision, boundary breach, physical-stop failure, or broken takeover chain	Safety red lines fixed; efficiency threshold pending preregistration; unexecuted



G1 first test: preregister, stop, recover, and independently retest

Innovation Is Not a Slogan: A Falsifiable Register

V2 closes another remaining gap: the proposal already introduces a proof line, a four-gate system, reversible public space, and public disclosure of failures, but “looking new” still does not prove that an innovation works. `visual/assets/innovation-register.json` binds six core innovations to their baseline shortcoming, novelty claim, falsifiable hypothesis, minimum evidence, pass threshold, failure signal, and exit action. INNOV-006 specifically tests whether the proposal creates a demonstrable and recoverable net increment over the opened park. A claim without a failure signal or exit action cannot be registered as a formal innovation, and operating results that do not yet exist remain unknown `[data:visual/assets/innovation-register.json#INNOV-001]` `[data:visual/assets/innovation-register.json#INNOV-006]`. Reviewers can therefore ask not only “what is new?” but also “what evidence would prove that it did not work?”

The evidence is organized into four tiers, each limited to a role consistent with its authority:

- Tier 1 comprises the open-call announcement, agent taskbook, and project site package, which define the task and submission boundary `[source:DATA-SRC-OFFICIAL-ANNOUNCEMENT-20260509]` `[source:AGENT-TASKBOOK]` `[source:SITE-PACKAGE]`.
- Tier 2 comprises the repository source registry, standards index, and processing guide, which locate evidence and its use limits `[source:SOURCE-REGISTRY]` `[source:PROCESSED-FACT-PACK]`.
- Tier 3 comprises public Beijing materials concerning the AI Origin Community, real-world testing along the century-old Jing-Zhang corridor, and the “one core, multiple points” innovation-district pattern. These materials assess possible coordination directions only; they do not mean that the project has been approved or that an organization has committed to participating `[source:BEIJING-AI-ORIGIN-2026]` `[source:BEIJING-AI-DISTRICTS-2026]`.
- Tier 4 combines national policies on data, AI-content labeling, and “AI+” with six global cases for mechanism inspiration and governance boundaries only `[source:NATIONAL-DATA-INFRA-2025]` `[source:AI-CONTENT-LABEL-2025]` `[source:AI-PLUS-2025]`.

Evidence Is Not a One-Time Snapshot: Expiry Must Propagate Downstream

`sources.json` records 29 sources and their access dates, but access on the same date does not mean that content remains valid indefinitely, nor does it reveal later revisions, replacement, unavailability, or the substitution of an official boundary for a provisional one. The new `visual/assets/evidence-freshness-policy.json` groups sources into project materials, provisional spatial data, urban context, policies and standards, and case references, then specifies review triggers, recommended maximum periods without re-verification, responsible roles, and invalidation actions. When a source becomes `review_due`, superseded, or unavailable, affected claims, matrix items, and scenario gate levels must be downgraded or frozen accordingly. The package currently confirms access dates only; it does not claim that a refresh audit with content summaries has been completed, so the completed-refresh count remains 0. Provisional spatial data remains `provisional_only` regardless of how recently it was accessed `[data:visual/assets/evidence-freshness-policy.json#FRESH-01]` `[data:visual/assets/evidence-freshness-policy.json#FRESH-02]`.

The professional response separates standards by question instead of stacking identifiers behind one conclusion:

- The open-call announcement and agent taskbook constrain the project and agent tasks `[standard:PROJECT-OFFICIAL-ANNOUNCEMENT]` `[standard:PROJECT-AGENT-OPEN-CALL-TASKBOOK]`.
- The urban-design scope and regulatory-planning process boundary are constrained by two housing and urban-rural development references `[standard:MOHURD-URBAN-DESIGN-MEASURES]` `[standard:MOHURD-CONTROL-DETAILED-PLANNING]`.
- Land-use terminology and the architectural-depth gap are recorded separately; a gap record is not represented as a verified clause `[standard:MNR-LAND-USE-CLASSIFICATION-GUIDE]` `[standard:MOHURD-ARCH-DESIGN-DEPTH-2016]`.

Precise red lines, existing buildings, ownership, regulatory-plan indicators, municipal infrastructure, traffic engineering, and heritage controls have not yet been obtained. `geometry/site_boundary.geojson` and the three key areas continue to use provisional repository extents for conceptual review only `[data:geometry/site_boundary.geojson#SITE-001]` `[source:DATA-SRC-PROVISIONAL-BOUNDARIES-20260605]`. They must not be used for approval, land acquisition, precise area calculation, or engineering implementation; boundary and control assumptions remain explicit `[assumption:A-BOUNDARY-001]` `[assumption:A-CONTROLS-001]`.

The latest boundary-basis note on the main branch also records an independent background check: OSM-mapped park features have 0% overlap with PROV-SITE-001, with a nearest distance of about 412.5 m. The reading may reflect incomplete OSM coverage, error in the provisional inferred extent, or both. It does not decide which geometry is correct and must not be used to alter or promote either one into an official boundary. This proposal therefore leaves the provisional geometry unchanged and treats receipt and verification of an official polygon, coordinate reference system, and version as a trigger for full recalculation `[source:PROVISIONAL-BOUNDARY-BASIS]`.

Three-Level Scope Framework

The three scope levels address different questions at different levels of precision while sharing one evidence chain. The approximately 43.6 km² coordinated research area asks how industry, research, urban problems, and external innovation nodes can collaborate. The approximately 11.4 km² overall design area asks how the century-old Jing-Zhang corridor can organize public space, slow mobility, functions, and scenarios. The three provisional key areas ask how three state stations—“co-create, verify, publish”—can shape blocks, buildings, public frontages, and operating gates. The areas express task hierarchy only and cannot be used to infer statutory boundaries `[metric:site_area_sqm]` `[metric:key_area_total_sqm]` `[depth:three_level_scope_framework]`.

The Single Core Mechanism

The century-old Jing-Zhang corridor is not a technology corridor on which AI products are placed. It is a public AI innovation production line that transforms urban problems into outcomes that are co-creatable, verifiable, pausable, reproducible, and deliverable to society. “AI Pilgrimage” does not mean worshipping technology. It describes a knowledge journey on which the public can “understand history → raise a problem → participate in joint testing → witness failure → leave a credited contribution.” Proof in the English name refers simultaneously to spatial evidence, technical verification, and proof of public value.

The overall structure evolves from “one spine, three anchors, six interfaces, and multiple nodes” to “one proof line, three state stations, two supply wings, and six urban interfaces.” The three state stations are not three homogeneous parks. Beijing AI Origin Community is responsible for co-creation and public deliberation; Zhongzhiyuan AI Independent Innovation Acceleration Area (Zhongzhiyuan) is responsible for technical, safety, and governance verification; and Dazhongsi is responsible for public release, urban services, and outcome transformation. The Xiaoyue River Scenario Enablement Wing supplies real problems and experiential feedback, while the Zhongguancun Technology Services Wing offers proposed support in intellectual property, compliance, talent, capital, and transformation. Every role assigned to an external organization is a planning recommendation and does not mean that cooperation has been secured `[assumption:A-EXTERNAL-COLLAB-005]` `[depth:overall_spatial_structure]`.


Readable Mapping to the Three Positioning Statements and Five Functions

The mapping below translates the taskbook’s own terms directly into spatial and operating mechanisms instead of leaving them as compliance-matrix checkmarks. It remains a conceptual response for professional teams to develop further; it creates no statutory function, project authorization, or institutional commitment [source:AGENT-TASKBOOK] [standard:PROJECT-AGENT-OPEN-CALL-TASKBOOK].

Taskbook position / function	Direct response in this proposal	Boundary that must not be crossed
The Centennial Jing-Zhang Cultural Belt	A double helix of railway engineering, Zhongguancun innovation, and AI correction history, with three knowledge-journey landmarks	Historical facts, heritage, names, and media await official/archival and rights review
The Urban AI Life Experience Belt	Six interfaces, nine everyday-first public nodes, non-AI parity, and quiet/screen-free modes	Field coverage is not an open service or a public-experience result
The AI Convergence Innovation Belt	A problem–co-create–verify–publish/service–transform–feed back state machine across three areas and two wings	No fabricated company, partner, land, funding, or implementation commitment
Full-stack independent AI innovation system	Zhongzhiyuan hosts proposed offline, shadow, bounded-test application, safety, and independent-retest gates	No claim of an existing full-stack platform, operator, or approved test field
World-class AI innovation ecosystem	Origin Community co-creation interfaces, six international comparisons, and two-wing factor support	“World-class” is an objective, not a ranking, concentration measure, or partnership fact
New paradigm for AI+ scenario enablement	Twelve scenario passports, three first-test protocols, and non-skippable G0–G3 gates	All items remain G0; approvals, participants, executions, and results remain 0
Intelligent AI-vibrant city	Daily, learning, Beta, quiet, and offline modes alongside staffed services	Screens, surveillance, event popularity, and mandatory accounts are not proxies for vitality
Global voice in AI governance	Traceable sources, published failures, suspension and grievance, expiry and retirement, and retest packages	A discussable method only; it is not an international rule or government institution

In conceptual south-to-north order, the spatial interfaces are I01 Dazhongsi Communication Interface, I02 Urban Services Interface, I03 Xiaoyue River Experience Interface, I04 University Co-creation Interface, I05 Zhongzhi Verification Interface, and I06 Qinghe Ecological Calibration Interface. They correspond to PUBLIC-004—PUBLIC-009 in geometry/public_space.geojson. Human-facing spatial prose uses I01—I06 only; G0—G3 is reserved for scenario maturity. The existing GeoJSON GATE-01—GATE-06 values are retained machine-compatibility identifiers, not spatial stages or human labels. All six interfaces are audit/service types pending field verification, not newly sited landmarks. Each interface type must jointly audit public space, a lateral slow-mobility link, sponge-system stitching, non-AI access, quiet hours, and scenario relay, prioritizing reuse of existing entrances and service facilities [data:geometry/public_space.geojson#PUBLIC-004] [data:visual/assets/site-grounding-register.json#SG-003] [metric:gateway_count].

The proposed logo uses two open-ended rail lines to form the letters JZ. The opening between them becomes a “verification opening with room for human intervention,” and six short ticks represent the six spatial interfaces. The logo uses only geometric linework and a project-owned wordmark; it does not use corporate trademarks, portraits, or restricted fonts. The overall palette comprises Rail Silver, Haidian Blue, Open-source Green, Verification Orange, and Bell Gold. Cultural wayfinding follows a separate three-line grammar—“mileage, year, source”—so that the cultural signage is not confused with the overall logo system.



Case	Transferable mechanism	Local Jing-Zhang action	Proposed verification	Not directly transferable
Barcelona 22@	Industrial upgrading proceeds in parallel with neighborhood renewal and knowledge transfer	Use the six spatial interfaces first to repair everyday public frontages	Whether public-interest delivery precedes expansion	Land-use policy and fiscal instruments [source:CASE-22AT]
King’s Cross	Heritage reuse, car-free public space, and mixed use	Present railway engineering history and contemporary public life along the same line	Whether residents continue to use the area on non-event days	Ownership, protection grades, and business model [source:CASE-KINGS-CROSS]
STATION F	Physical space operated through programs, mentors, and partner networks	Four-season operating protocol and a problem-to-Proof funnel	Cycle time from problem to bounded prototype	A single large campus and private operating conditions [source:CASE-STATION-F]
MaRS	Coupling laboratories, offices, events, and community services	Co-locate publication, compliance advice, and staffed service at Dazhongsi	Success rate of human handoff and follow-up services	Healthcare, investment, and institutional systems [source:CASE-MARS]

Regional coordination uses a “testing relay,” not a vague alliance. The Future Science City is proposed for retesting in basic science or frontier equipment; Huairou Science City for verification related to major scientific facilities; the Beijing Economic-Technological Development Area for engineering retests in manufacturing and embodied systems; other first-batch innovation districts for cultural or industrial scenarios in which they specialize; and Beijing–Tianjin–Hebei nodes for cross-regional replicability checks. Each relay transmits only tasks, protocols, anonymized results, and failure records; it does not require the transfer of raw sensitive data. The publicly stated “one core, multiple points” pattern only shows that this division of work has a basis for discussion; it does not mean that any node has agreed to participate [source:BEIJING-AI-DISTRICTS-2026].


Overall Design Area: Urban Renewal and Regulatory-Plan-Level Urban Design

The overall area takes the opened park’s north–south green corridor, reported lateral connections, and all-age public activity as a publicly reported baseline rather than an object that this proposal has yet to build from scratch [source:HAIDIAN-JZ-PHASE2-OPEN-2026]. Only then does it organize a three-layer section: railway-side historical interpretation, existing slow-mobility and blue-green space, and a reversible innovation-service layer. The railway side receives only low-intervention interpretation; the middle first protects everyday walking, cycling, and staying; and the neighborhood side uses six interface types to audit breaks and add small-scale services. Any action crossing railways, rivers, urban roads, bridges, tunnels, or underground space is only a question pending verification and a design direction; separate field, transport, flood-control, railway-safety, and engineering studies are mandatory.

Land use is refined from 12 large color blocks into 36 irregular, topologically closed conceptual units. They cover nine project-enumerated categories: research, education, housing, community services, commercial services, culture, green space, plazas, and flexible reserve. The number and categories of units can be recalculated from the GeoJSON [metric:land_use_unit_count] [metric:land_use_category_count]. They express spatial–industrial functional envelopes and are neither statutory parcels nor approved uses [data:geometry/land_use.geojson#LU-001] [depth:land_use_layout].

The six lateral interface lines and the north–south main line form a fishbone accessibility network. Every interface binds slow mobility, public space, stormwater stitching, and a testable urban problem. The nine public spaces do not turn the park into a round-the-clock testing ground. Instead, they follow reversible time-space rules: daily life is the default; public learning is scheduled; bounded Beta operation occurs only during approved periods; quiet mode applies from 22:00 to 07:00; and everyone can always choose a non-AI route. The innovation is not that “the city adapts to AI,” but that “AI must adapt to everyday urban life.”

The building layer uses massing prototypes to express priority intervention capacity; it does not simulate the complete existing city. Each prototype is marked as retain for assessment, adaptive-reuse candidate, or new-build candidate. Without evidence from surveys, ownership, structural conditions, heritage controls, and regulatory planning, none may be converted into a demolition or construction conclusion. Floor area ratio, total floor area, building height, road red lines, parking, and municipal capacity remain unknown [assumption:A-SPATIAL-REPRESENTATION-002] [depth:development_intensity_controls] [depth:height_massing_character].



Land-use structure and spatial transmission

Detailed Design of Key Areas

The three key areas use the same response framework—site conflict, spatial structure, core project, AI scenario, admission gate, and acceptance—and organize the material at key-area detailed-design depth [depth:three_key_area_detailed_design]. The extents below remain provisional substitutes; all area and massing-coverage values are for conceptual review only [data:geometry/key_areas.geojson#PROV-KEY-001] [data:geometry/key_areas.geojson#PROV-KEY-002] [data:geometry/key_areas.geojson#PROV-KEY-003].

Key-area evidence crosswalk

To keep “key area → scenario → acceptance” reviewable rather than implicit in long prose, visual/assets/key-area-evidence-matrix.json crosswalks each provisional area to its state station, public anchor, I/GATE interface references, AI service zone, scenario nodes and concept-level acceptance contract, then lists the formal evidence required before G1 [data:visual/assets/key-area-evidence-matrix.json] [metric:key_area_evidence_matrix_record_count]. Its 100% matrix-field coverage means only that the three records contain the required crosswalk fields; it does not mean field coverage, partner confirmation, approval, operating performance or public-value evidence. Field audits, accountable-role confirmation, approvals, test executions and known results remain 0 for all three records [metric:key_area_evidence_matrix_field_coverage_ratio].

Provisional key area	Station / conceptual spatial response	AI zone / scenarios	Evidence required before G1	Current boundary
PROV-KEY-001 Zhongzh yuan	VERIFY; one garden, one stack, two interfaces	AI-ZONE-001; SCENE-001—004	Legal/owner site, safety and physical-stop review, accountable roles, non-AI baseline, independent retest package	Provisional extent, not approved, untested
PROV-KEY-002 Origin Community	CO-CREATE; one street, two courtyards, four nodes	AI-ZONE-002; SCENE-005—008	Venue and safeguarding, rights and consent, non-technical participation, withdrawal route, independent retest package	Provisional extent, no partner commitment, untested

Provisional key area	Station / conceptual spatial response	AI zone / scenarios	Evidence required before G1	Current boundary
PROV-KEY-003 Dazhongsi	PUBLISH; four-quadrant walking plus one commons and one desk	AI-ZONE-003; SCENE-009—012	Route/service ownership, staffed desk, source maintenance, same-task comparator, independent retest package	Provisional extent, no approved service desk, untested

The matrix adds only reviewable design mappings. It does not turn interface IDs into sited landmarks or public desk research into as-built survey, regulatory plan, ownership or institutional commitment evidence [data:visual/assets/key-area-evidence-matrix.json].

Zhongzhiyuan: Verification Station / VERIFY

The conflict is that autonomous technology needs real problems and compound testing, while urban public space cannot become an unbounded testing ground. In line with Haidian’s latest public objectives, Zhongzhiyuan treats AI4S, AI-safety governance, and mutual recognition of rules as proposed verification themes; it does not claim that platforms, partners, or international mechanisms already exist [source:HAIDIAN-JZ-MIDTERM-2026]. The spatial structure is “one garden, one stack, two interfaces.” Qinghe Test Garden supports environmental and low-risk joint testing. The Verifiable Model Commons supports offline benchmarks, safety checks, and interoperability checks. The community-facing side provides a public observation and grievance interface, while the service side supports human takeover and equipment maintenance. AI-ZONE-001 binds four nodes: low-speed delivery, energy use, safety assessment, and the failure archive [data:geometry/constraints.geojson#AI-ZONE-001] [data:geometry/constraints.geojson#SCENE-001]. Admission requires complete sources, a responsible person, a risk level, a comparison baseline, and a physical emergency stop. Any collision, boundary breach, unauthorized data processing, or failed human takeover triggers immediate suspension. Acceptance is not based on a “successful demonstration,” but on a complete retest package, reproducible failure, usable human takeover, and passage through the public-value gate.


Beijing AI Origin Community: Co-creation Station / CO-CREATE

The conflict is the absence of a stable interface between a high density of innovation resources and residents’ daily lives, low-cost entrepreneurship, and public participation. Public materials describe Origin Community as an approximately 3 km² area with a “ten-minute innovation circle” context; published scale figures are context only, not an inventory of this project’s resources, a siting basis, or evidence of partnership [source:HAIDIAN-AI-TRIAL-FIELD-2026]. The spatial structure is “one street, two courtyards, four nodes,” with five-minute everyday-life support cells nested inside the ten-minute innovation context and a broader fifteen-minute public-service planning framework. The Open-source Release Street connects universities and the community. The Co-creation Courtyard supports problem decomposition and team matching. The Public Deliberation Courtyard supports resident joint testing, education for minors, and screen-free discussion. AI-ZONE-002 binds trusted cultural guidance, open-source matching, educational workshops, and climate-and-accessibility joint testing [data:geometry/constraints.geojson#AI-ZONE-002] [data:geometry/constraints.geojson#SCENE-005]. Admission requires a problem grounded in an explicit public need, participants able to withdraw, and recommendations excluded from hiring evaluation. Acceptance focuses on cross-role co-creation, responsiveness to withdrawal, non-technical participation, and differences in task success across groups.

Dazhongsi: Publication Station / PUBLISH

The conflict is that station-area footfall, enterprise services, and consumer display can easily turn AI into screen-based marketing while displacing staffed services and quiet space. The design first acknowledges the opened southern section’s community, commuting, and all-age activity baseline, rather than repackaging existing connectivity improvements as this proposal’s outcome [source:HAIDIAN-JZ-PHASE2-OPEN-2026]. The spatial structure is “four-quadrant walking + one hall and one desk.” The four quadrants indicate connection directions pending field verification. The Bell-and-Rail Commons presents evidence of passage, failure, and correction. The Urban Services Desk places intelligent navigation alongside a staffed counter. The publication frontage provides screen-free default areas and quiet hours. Haidian’s public Fifteenth Five-Year Plan sets district-level objectives to become a global AI innovation source and industry highland by 2030 and to advance high-quality urban renewal; it is used here only as district direction and does not show that any specific Dazhongsi parcel is included, rights are settled, or delivery is committed [source:HAIDIAN-15FYP-2026]. AI-ZONE-003 binds accessible routing, slow-mobility guidance, enterprise services, and health navigation [data:geometry/constraints.geojson#AI-ZONE-003] [data:geometry/constraints.geojson#SCENE-009]. Acceptance focuses on source-hit rate, expiration notices, human handoff, complaint closure, and incremental value over existing services—not event attendance or exposure.

The three pilgrimage landmarks are redefined as three chapters in one knowledge journey. Zhongzhiyuan’s “Open-source Spark Tower” records reproducible contributions without displaying personal rankings. Origin Community’s “Algorithm Milestone” places Jing-Zhang engineering history, Zhongguancun innovation history, and the history of AI correction side by side. Dazhongsi’s “Bell-and-Rail Commons” houses an archive of failures and corrections. Contributors may choose real names, pseudonyms, or anonymity. Recognition records only verifiable public contributions and is not tied to wealth, traffic, recruitment, or administrative evaluation. All names and forms require copyright, trademark, portrait-rights, historical-source, and accessibility review [assumption:A-CULTURE-CONTENT-006].



Spatial refinement of the three state stations

AI Innovation Ecosystem, Personas, and AI+ Scenarios

Seven Joint-Testing Profiles

Haidian’s latest public requirements prioritize three groups: leading research talent, young innovators and entrepreneurs, and permanent residents. This proposal does not create a separate persona system. Instead, it maps them to the existing seven: leading research talent primarily maps to open-source developers and university faculty and students; young innovators and entrepreneurs map to start-up teams and connect with enterprise-service teams; permanent residents map to nearby residents and older people and people with reduced mobility; tourists and international visitors remain as reviewers of communication and legibility. The three policy groups set demand priorities, while the seven profiles decompose joint-testing and risk. Neither means that samples have been recruited or research completed [source:HAIDIAN-JZ-MIDTERM-2026] [data:visual/assets/site-grounding-register.json#SG-006].

Profile	Core needs	Potential harm	Hard protection
Open-source developers	Computing, protocols, retesting, and collaborators	Contributions captured by a platform	Right to withdraw, clear licenses, portable contributions
Start-up teams	Low-cost verification, compliance, and transformation	Misled by demonstration metrics	Publish failures; prohibit implied government endorsement
Enterprise-servic e teams	Clearly sourced administrative and market information	Harm from outdated content	Source dates, human handoff, and clear responsibility boundaries

Profile	Core needs	Potential harm	Hard protection
University faculty and students	Learning, research, and outcome transformation	Leakage of education data	Local sandbox, teacher review, and minimum data for minors
Nearby residents	Quiet, accessibility, and everyday services	Noise, congestion, surveillance, and exclusion	Daily life first; quiet and screen-free operation; grievance and pause rights
Older people and people with reduced mobility	Continuous accessibility and human assistance	Exclusion through digital barriers	Non-AI access, joint testing, and no withdrawal of staffed service
Tourists and international visitors	Trustworthy multilingual cultural content	False narratives and excessive collection	Source labels, curatorial review, and account-free browsing

Non-AI Options Must Complete the Same Basic Task

Existing indicators can show that nine public nodes declare non-AI access and that twelve scenarios declare manual fallback, but they cannot prove that those alternatives are actually usable. The new `visual/assets/non-ai-parity-contract.json` requires every public service seeking entry to G2 or G3 to compare six dimensions: opening hours, basic outcome, cost, rights and grievance, safety and accessibility, and completion time. A person who does not use an account, smartphone, or algorithm must not receive lower eligibility, higher charges, or weaker review rights. No on-site joint-test results currently exist, so all four test journeys remain unknown; any permitted completion-time difference must be preregistered after a baseline is obtained rather than invented in the proposal `[data:visual/assets/non-ai-parity-contract.json#PARITY-001]` `[data:visual/assets/non-ai-parity-contract.json#PARITY-004]`.

Four-Gate System and Scenario Passports

JZ-AIOS requires every scenario to proceed in sequence through G0 concept/offline, G1 shadow comparison, G2 time-bounded and geographically bounded Beta, and G3 approved public operation. No level may be skipped. Every node in the current submission is in G0 conceptual status; no scenario is operating or approved. Each passport must include version, responsible role, risk level, minimum dataset, geographic and temporal boundary, potentially harmed groups, human takeover, grievance channel, non-AI option, admission conditions, stop conditions, retest package, and expiry date `[assumption:A-SCENARIO-PROTOCOL-004]`.

The 12 scenarios are grouped in human-facing prose as safety and access (SCENE-001, 003, 009, 010), trusted public service (SCENE-002, 005, 011, 012), and public learning and collaboration (SCENE-004, 006, 007, 008). The machine record retains every SCENE-001—012 passport and field; grouping does not indicate deployment, priority, or test results. The ten service scenario cards are listed below. A spatial Feature indicates conceptual placement only, not deployment:


ID / Node	Scenario and space	Minimum input → output	Responsibility and manual fallback	Admission / stop / retirement
SC-01 / SCENE-005	Trusted Jing-Zhang cultural guide; Origin-milestone route	Rights-cleared historical sources and an actively selected language → sourced content	Human curator; account-free print/audio alternative	Complete sources and AI labels; remove disputed content that cannot be verified; annual review
SC-02 / SCENE-009	Accessible route assistant; Dazhongsi and I01—I06	Slopes, temporary obstacles, and facility status → suggested route	Confirmation by on-site service staff; retain physical wayfinding	Road test and joint testing passed; stop recommending a route upon any critical obstacle; invalidate when data changes
SC-03 / SCENE-010	Slow-mobility crowding notice; public greenway	Segment-level anonymous flow → detour suggestion	Published by management staff; no automated control	No facial data; return to static wayfinding if false alerts affect passage; expire after the event
SC-04 / SCENE-001	Low-speed delivery sandbox; authorized surface in Zhongzhiyuan	Vehicle state, geofence, and obstacles → low-speed task	On-site safety officer, physical emergency stop, and manual towing	May apply for G2 only after passing offline and closed-site tests; one collision or boundary breach triggers suspension; retire the configuration after each test
SC-05 / SCENE-011	Enterprise-service agent; Dazhongsi Services Desk	Public policy and a user-initiated question → sourced navigation	Staff handle formal matters	Complete sources and update dates; suspend if an error affects official procedures; policy updates trigger invalidation
SC-06 / SCENE-006	Open-source collaboration matching; Origin Co-creation Courtyard	Voluntarily disclosed skills and topics → team suggestions	Community operator processes withdrawal; excluded from hiring evaluation	Confirm withdrawal and license terms; stop upon discriminatory matching; clear data when the project ends
SC-07 / SCENE-002	Public-space energy assistant; Zhongzhiyuan and I01—I06	Equipment and environmental data → lighting suggestion	Operations staff retain priority; manual switches remain available	No identity data; return to manual control if safety lighting is affected; retest when seasonal parameters expire
SC-08 / SCENE-012	Community health navigation; beside staffed service at Dazhongsi	Publicly listed institutions and a user-initiated inquiry → public information and booking entrance	No diagnosis; emergencies handed off to professional institutions	Clear boundary notice; any diagnostic output triggers suspension; invalidate when institution information changes
SC-09 / SCENE-007	Educational experience workshop; Origin Community	Local teaching data → bias and programming exercises	Full teacher review; minors’ data do not leave the site	Parent/school rules and local sandbox in place; stop upon external data transfer; destroy temporary data after the course
SC-10 / SCENE-003	Assisted safety-incident assessment; Zhongzhiyuan control room	Equipment alerts and human reports → unverified incident summary	Authorized personnel decide; no automated enforcement	Permissions, logs, and comparison baseline complete; suspend upon overreach or misleading handling; re-review after model changes


Two public-governance nodes are added. SCENE-004, the “Public Archive of Failures and Corrections,” records suspensions, corrections, and voluntary retirements. SCENE-008, “Climate and Accessibility Joint Testing,” makes older people, children, and people with reduced mobility co-testers at an early stage rather than passive recipients of acceptance after completion. Node and service-zone counts are machine-verifiable `[metric:mapped_scenario_node_count]` `[metric:ai_service_zone_count]`. Passport and manual-fallback field coverage is counted separately and does not prove field performance `[metric:scenario_passport_coverage_ratio]` `[metric>manual_fallback_coverage_ratio]`.

Three Priority Industrial Test Protocols

- T-01 Robot Safety Graduation Protocol: G0 digital and closed-site baseline → G1 shadow task → application for G2. Proposed hard gates require all physical emergency stops to succeed in controlled tests, zero collisions, zero boundary breaches, and a complete human-takeover chain. Any safety red line triggers suspension, review, and restart from a lower gate. The figures are proposed design acceptance criteria, not measured results.
- T-02 Sourced Enterprise-Service Protocol: This is the first offline reproducible evidence package: it freezes the question set, sources, update times, refusals, and human handoff so others can review the same materials. It is a conceptual test package, not an approved test or field result. `visual/assets/g0-offline-enterprise-service-baseline.json` fixes five questions: source status, opening background, policy direction, provisional-boundary refusal, and human handoff for formal procedures; completed replays, answer outputs, and field tests are all 0. Its 5/5 field coverage is protocol completeness only, not service accuracy `[data:visual/assets/g0-offline-enterprise-service-baseline.json]` `[metric:g0_offline_contract_coverage_ratio]` `[metric:completed_g0_offline_replay_count]`. A bounded service should be entered only when source coverage, stale-information warnings, and human handoff all reach pre-published thresholds. Suspend the service if an error affects a formal procedure.
- T-03 Universal-Access Route Protocol: This is the spatial and public-interest exemplar: wheelchair users, visually impaired users, older people, child caregivers, and ordinary pedestrians complete the same task together. A proposal cannot graduate if any critical break lacks a non-AI alternative route. It likewise has no approval or field result; only after operation begins may it continue comparing gaps in task success among groups and publish improvements.

Graduation uses a three-dimensional cube of “technical maturity × public-value maturity × governance maturity.” Failure in any dimension allows only correction or retirement. The national “AI+” policy emphasizes openness and sharing, safety and controllability, and broadly shared outcomes. This proposal translates those directions into a public-value gate but does not replace specific law, approvals, or professional responsibility `[source:AI-PLUS-2025]` `[depth:municipal_new_infrastructure]`.





Six spatial interfaces, slow mobility, blue-green systems, and scenario nodes

Blue-Green Network, Public Space, and Urban Character

The blue-green system no longer claims to build the main green spine from scratch. The publicly reported approximately 9 km green corridor, north–south sections, fishbone slow-mobility network, and sponge facilities become an existing baseline pending field verification. The proposal’s increment is limited to continuity audits at the six interface types, low-risk learning interfaces at the three state stations, and reversible components that do not damage everyday recreation [source:HAIDIAN-JZ-PHASE2-OPEN-2026] [data:visual/assets/site-grounding-register.json#SG-004]. Green-space area and count are still recalculated from submission-package design geometry [metric:green_space_area_sqm] [metric:green_space_count] and cannot be represented as as-built conditions. Because the ratio uses the provisional boundary as its denominator, it inherits low confidence [metric:green_ratio] [depth:blue_green_public_space].

The nine public spaces consist of three pilgrimage-landmark living rooms and six spatial interfaces, forming a line–surface–point network. Daily-life mode is the default at every space. Public learning or bounded Beta begins only when notice, authorization, time limits, and rollback conditions are all in place. Quiet mode applies from 22:00 to 07:00. Screen-free areas are the default, and bright large displays, continuous broadcasting, and dense camera arrays are not treated as technological character. Public-node count and area can be recalculated [metric:public_space_node_count] [metric:public_space_area_sqm]. The ratio remains a low-confidence conceptual value; screen-free defaults and the spatial layer retain separate review entry points [metric:public_space_ratio] [metric:screen_free_public_node_ratio] [data:geometry/public_space.geojson#PUBLIC-001].

The public-space component library contains six replaceable component types: bilingual/tactile milestones; offline maps and human-call points; switchable low-level guide lighting; mobile joint-testing tables; rain gardens and shaded seating; and bounded test barriers with physical emergency stops. Every component must be usable without an account, retain basic function when power is lost, and identify who is responsible for maintenance. Bridges, tunnels, equipment foundations, and utility conditions require separate investigation.

Culture follows a “double helix of trustworthy narrative.” The physical strand explains railway engineering → urban growth → Zhongguancun innovation. The learning strand explains problem raising → model failure → human correction → public verification. Public Beijing materials show that Phase I of the Jing-Zhang Railway Heritage Park connects universities, research, and residents’ lives through railway remains, green public space, and functional stitching. Public participation and the responsible-planner mechanism in the existing project also show that cultural space requires continuous professional stewardship [source:JZ-PARK-2023] [source:JZ-COCREATION-2021]. V2 does not copy the existing landscape; it turns “source, correction, and failure” into new cultural content.

All AI-generated text, images, audio, video, or virtual scenes carry both public-facing explicit labels and, where applicable, file-metadata labels. Disputed content enters human curatorial review and is removed if it cannot be verified. This rule is consistent with the direction of the 2025 synthetic-content labeling regime, but formal operation still requires legal review [source:AI-CONTENT-LABEL-2025]. The international message is “Build in public. Test with people. Keep the evidence.” It emphasizes open development, joint testing, and retained evidence rather than claiming global leadership or completed delivery.

Renewal Projects, Implementation Policy, and Phasing

This section places phases, pilots, participating parties, and measurable indicators within one implementation framework. Each project first states prerequisites and then a proposed responsibility combination and acceptance indicators. At every phase, monitoring, resident feedback, and professional assessment must determine whether to continue, correct, or exit.

Eight projects form a renewal package that maps space, protocols, and operations one to one [depth:renewal_project_list]:

Project	Main content and Feature	Prerequisites	Proposed responsibility combination	Conceptual acceptance
JZ-01 Accessible Stitching at Six Interfaces	PUBLIC-004—009 and the lateral interface lines	Existing-condition survey and transport/accessibility review	Public bodies + transport/landscape professionals + community	I01—I06 each have a continuous non-AI line, break-point register, and rollback plan
JZ-02 Reversible Time-Space Park	Nine public spaces, the main green spine, and sponge stitches	Green-space, water, railway, and operations conditions	Landscape/ecology professionals + community operator	Daily, learning, Beta, quiet, and offline modes can switch
JZ-03 Zhongzhi Verification Commons	AI-ZONE-001 and four nodes	Safety, network, fire safety, and responsible entity	Technical testing + safety/ethics + on-site operations	Retestable, emergency-stoppable, and able to publish failures
JZ-04 Origin Open-source Co-creation Street	AI-ZONE-002 and four nodes	Community co-creation, education, and copyright rules	Universities/developers + community + curation	Withdrawal, licensing, non-technical participation, and joint testing are verifiable
JZ-05 Dazhongsi Urban Services Portal	AI-ZONE-003 and the staffed service desk	Station-city, transport, commercial, and public-service conditions	Enterprise service + staffed counter + operations	Sources, expiry warnings, human handoff, quiet operation, and screen-free space all coexist
JZ-06 Trustworthy Narrative Double Helix	Three landmarks, milestones, and failure archive	Historical-source, copyright, trademark, and accessibility review	Curation + heritage/history professionals + public	Full coverage of sources, AI labels, correction, and alternative media

Project	Main content and Feature	Prerequisites	Proposed responsibility combination	Conceptual acceptance
JZ-07 Scenario Passports and Open Benchmarks	JZ-AIOS, retest packages, and annual expiry	Legal, safety, data, and professional review	Independent evaluation + responsible operations + community seats	Passport fields, manual fallbacks, and stop thresholds are complete
JZ-08 Four-Season Operations and Proof Week	Annual problems, open source, Beta, and evidence release	Budget, venue, responsible party, and event approval	Proposed Jing-Zhang Open Collaboration Desk	Publish passes, failures, corrections, and voluntary retirements together

visual/assets/pilot-readiness-register.json adds conceptual RACI, approval triggers, prohibited data, incident and shutdown responsibility, community co-testing, exit and recovery, independent retesting, and go/no-go evidence for JZ-01—JZ-08 and T-01—T-03. “Coverage” means only that all eight projects and three pilots have these fields; it does not mean that responsible parties have accepted roles, approvals have been obtained, or pilots have operated. Current status remains uniformly G0/concept_only, while field outcomes and recovery time remain unknown [data:visual/assets/pilot-readiness-register.json#JZ-01] [metric:pilot_readiness_protocol_coverage_ratio].

Phasing uses “year window + evidence gate”; a date never automatically grants eligibility to enter the next phase [data:geometry/phasing.geojson#PHASE-1] [metric:phase_count] [depth:phasing_implementation]:

Window	Status	Admission gate	Graduation gate	Rollback
Years 0—2	G0 concept/offline and G1 shadow	Complete sources, responsibility, risk, baseline, and manual fallback	Safe rollback, public-value gate passed, and complete retest package	Any major safety, privacy, accessibility, or public-value failure
Years 3—5	Approved G2 bounded Beta	Independent retest, community deliberation, professional review, and site permission	Pass the three-dimensional graduation cube and publish evidence	Repeated failure, unanswered grievances, or undelivered public benefit
After Year 5	Apply for G3, expand, correct, or retire	Formal approval and operating resources defined	Annual report confirms continuation, correction, or voluntary retirement	Major change in boundaries, controls, technology, or social impact

Annual operations follow a four-season protocol. In the first quarter, “Problem Season,” communities, universities, and urban service providers publish problems suitable for public release. In the second quarter, “Open-source Season,” teams, protocols, and offline baselines are formed. In the third quarter, “Urban Beta Season,” only approved bounded tests operate. In the fourth quarter, “Proof Week,” passes, failures, complaints, corrections, and voluntary retirements are published, generating the following year’s problems. The developer-transformation funnel is “public problem → co-creation team → offline baseline → bounded application → public evidence → staffed service / enterprise transformation / external retest.” Participation must never be presented as a commitment to investment promotion, funding, or policy. Event branding uses seasonal colors from the same logo and does not create a separate IP that conflicts with the belt’s primary brand.

Operations propose a “Jing-Zhang Open Collaboration Desk.” Public bodies define rules only within their statutory duties. Professional organizations review space, data, safety, and accessibility. Universities and firms supply withdrawable projects. Community representatives hold agenda seats and rights to grievance and proposed suspension. The annual report must disclose services, public benefits, complaints, data incidents, human takeovers, energy use, failures, and improvements together. Target values must not be presented as achievements when measured data do not exist.

Metrics, Area Recalculation, and Compliance Matrix

The seven indicator categories are presented as readable review questions:

- Site and key-area scale: provisional site area, key-area count, and total key-area area [metric:site_area_sqm] [metric:key_area_count] [metric:key_area_total_sqm].
- Land-use structure: land-use unit and category counts [metric:land_use_unit_count] [metric:land_use_category_count].
- Building representation: massing-prototype count, footprint, and overall representation ratio [metric:building_count] [metric:building_footprint_area_sqm] [metric:building_representation_ratio].
- Key-area massing: key-area prototype count and footprint coverage [metric:key_area_building_count] [metric:key_area_building_footprint_ratio].
- Blue-green system: green-space area, count, and ratio [metric:green_space_area_sqm] [metric:green_space_count] [metric:green_ratio].
- Public space: public-space area, node count, and ratio [metric:public_space_area_sqm] [metric:public_space_node_count] [metric:public_space_ratio].
- Transport: design-route count and length [metric:design_road_count] [metric:design_road_length_m].

The required area families now go beyond unit counts and are complete by recomputable object while retaining the conceptual-design boundary:

- The commercial-service, residential, and community-service envelopes are recalculated from the land-use layer; they are not statutory parcels or land-supply quantities [metric:land_use_area_05_sqm] [metric:land_use_area_0701_sqm] [metric:land_use_area_0702_sqm].
- Research and development, culture, and education/research envelopes likewise describe only the current design geometry; they create no floor-area ratio or construction quantum [metric:land_use_area_0802_sqm] [metric:land_use_area_0803_sqm] [metric:land_use_area_0804_sqm].
- Green-space, plaza/public-interface, and flexible-reserve envelope areas are recalculated separately; the 1403 functional envelope is not the same as the nine direct PUBLIC_SPACE intervention footprints [metric:land_use_area_1401_sqm] [metric:land_use_area_1403_sqm] [metric:land_use_area_16_sqm].
- Areas for the three year windows come from geometry/phasing.geojson. They show evidence-gate coverage only, not approved construction boundaries or an automatic delivery sequence [metric:phase_1_area_sqm] [metric:phase_2_area_sqm] [metric:phase_3_area_sqm].
- Each key area retains a low-confidence area recalculated from its provisional polygon. Published approximate areas and provisional-geometry areas remain separate fields, while official polygon areas remain unavailable [metric:key_area_zhongzhiyuan_sqm] [metric:key_area_origin_community_sqm] [metric:key_area_dazhongsi_sqm].

AI and public-value indicators are separated in the same way:

- Scenario scale: service-scenario, mapped-node, and service-zone counts [metric:scenario_count] [metric:mapped_scenario_node_count] [metric:ai_service_zone_count].
- Operating interfaces: spatial interfaces I01—I06, passport-field coverage, and manual-fallback-field coverage [metric:gateway_count] [metric:scenario_passport_coverage_ratio] [metric>manual_fallback_coverage_ratio].
- Public access: non-AI access and screen-free-default coverage [metric:non_ai_access_coverage_ratio] [metric:screen_free_public_node_ratio].
- Implementation: three phase gates [metric:phase_count].
- Field grounding and handoff readiness: site-grounding observations, responsibility-field coverage for eight projects and three pilots, and the post-opening field-audit count that remains 0 [metric:site_grounding_observation_count] [metric:pilot_readiness_project_coverage_ratio] [metric:completed_post_opening_field_audit_count].
- First-test documentation: all 12 scenes have preregistration records with 100% required-field coverage [metric:g1_preregistration_record_count] [metric:g1_preregistration_required_field_coverage_ratio].
- First-test real-world evidence: completed preregistrations and approved windows both remain 0 [metric:completed_g1_preregistration_count] [metric:approved_g1_test_window_count]; field executions and known independently retested results also remain 0 [metric:g1_field_execution_count] [metric:known_g1_result_count].

These indicators prove only “what has been designed” in the submission package; they do not prove real-world operating performance.

All area ratios using the provisional boundary as their denominator have low confidence. Counts of design objects and attribute-completeness rates may have high confidence. Route lengths and massing areas have medium conceptual-design confidence or low confidence where affected by the boundary. Floor area ratio, total floor area, building density, average height, road area and road ratio, parking supply, measured recovery time, and energy per effective service remain pending until statutory material or controlled-test evidence is available; massing-prototype coverage and road centerlines are not substitutes [metric:building_density] [metric:road_ratio]. The fixed evidence chain is “public source / explicit assumption → GeoJSON → EPSG:4548 recalculation or attribute count → metrics.json → text / drawings / HTML → machine self-check → human professional judgment” [depth:metrics_recalculation].

compliance_matrix.json no longer makes 23 tasks share the same evidence bundle. agent.1 points to branding and the three-zone/two-wing structure; agent.2 to the cases and industrial state machine; agent.3 to scenario passports and test protocols; agent.4 to nine public spaces, three landmarks, and the component library; agent.5 to trustworthy narrative and synthetic-content labeling; and agent.6 to four-season operations and the transformation funnel. Announcement tasks are also mapped separately to the relevant chapter, layer, metric, source, assumption, and self-check. standard_matrix.json and design_depth_matrix.json likewise assign genuine evidence by professional question instead of using a batch-copied summary.



From field coverage to evidence maturity

Risk, Copyright, and Compliance

- Boundary and planning risk: Official GIS/CAD, ownership, regulatory planning, road red lines, and engineering conditions are absent. Every spatial conclusion is only a conceptual recommendation. When formal information becomes available, it must replace the provisional data, be compared for differences, and trigger full recalculation.
- Spatial-representation risk: The 36 units and massing prototypes are design representations, not an existing-condition survey. The representation ratio must not be interpreted as building density or construction scale [assumption:A-SPATIAL-REPRESENTATION-002].
- Public-interest risk: Testing may create noise, congestion, digital exclusion, or privatization of public space. Everyday life, non-AI access, quiet and screen-free operation, joint testing, grievance, and suspension are hard gates.
- AI and data risk: Bias, hallucination, privacy leakage, unauthorized automation, and vendor lock-in are controlled through data minimization, retention at source, logging, human responsibility, retesting, expiry, and retirement. Automated enforcement, diagnosis, and substitution for formal approval are prohibited.
- Safety and resilience risk: Robot, vehicle, and equipment testing occurs only within authorized areas. Physical emergency stops, on-site safety officers, offline human operation, and L0/L1/L2 degradation must remain available. Energy and recovery indicators remain unknown until measured.
- Cultural and historical risk: Historical facts, people, artifacts, and engineering materials require verification through official, archival, or rights-cleared sources. AI-generated content receives explicit and metadata labels; disputed content can be corrected, removed, and traced [assumption:A-CULTURE-CONTENT-006].
- Copyright and branding risk: The file-level rights-status ledger covers all 54 manifest paths (including 53 non-manifest content files) and distinguishes self-declared originals, repository-provisional derivatives, and generated outputs pending audit. Completed independent file-level clearance audits remain at 0 and the overall status is not_fully_cleared. The complete COMMUNITY-DISPLAY-ONLY terms, OSM ODbL obligations, PDF fonts, generation-tool terms, editable sources, and logo/landmark trademark status still require review; the package must not claim full clearance [data:visual/assets/rights-clearance-ledger.json#RIGHTS-01] [data:visual/assets/rights-clearance-ledger.json#RIGHTS-GATE-02].
- External-coordination risk: Future Science City, Huairou Science City, the Beijing Economic-Technological Development Area, other innovation districts, and Beijing–Tianjin–Hebei are only optional retest roles. Without written confirmation, none may be described as a partner, investor, or committed implementation party [assumption:A-EXTERNAL-COLLAB-005].
- Operations and equity risk: Event popularity cannot replace resident satisfaction, accessibility, fairness, or complaint closure. Recognition for contributions must not be used for traffic rankings, employment screening, or administrative evaluation.
- Tool and evidence risk: Machine checks verify only structure, topology, references, and consistency. They do not replace professional judgment in planning, architecture, transport, municipal engineering, landscape, ecology, fire safety, railway safety, data security, accessibility, community engagement, or law [depth:risk_missing_data].
- Public-reporting versus field-condition risk: This Firecrawl desk-research pass preserves public-page sources, dates, summaries, and content digests for citation and design judgment only. It did not conduct a site visit, review as-built drawings, or audit facility operations. Every claim involving exact location, built condition, intensity of use, or accessibility performance requires field verification before G1 [data:visual/assets/site-grounding-register.json#SG-001].

All current AI scenarios are in G0 conceptual status. The eight projects and four-season events are proposals: they are not approved, built, or operating, and no organization has committed to them. Entry to a higher operating gate may be discussed only after statutory approval, responsible entities, professional review, public participation, funding and operations, and incident response are all defined.

References

The primary project basis comprises [source:DATA-SRC-OFFICIAL-ANNOUNCEMENT-20260509], the Beijing Municipal Commission of Planning and Natural Resources open-call announcement; [source:AGENT-TASKBOOK], the repository agent taskbook; [source:SITE-PACKAGE], the site package; [source:SOURCE-REGISTRY], the source registry; and [source:PROCESSED-FACT-PACK], the processing guide. The provisional spatial basis is [source:DATA-SRC-PROVISIONAL-BOUNDARIES-20260605], used only for generation, presentation, and informal review.

The direct entry to the repository's public-information index is brief/public-brief.md; public boundary and use limitations are described in brief/README.md. Both provide only public task context and data boundaries and cannot generate statutory control values.

Local professional materials include brief/site-package/standards/standards.json and its references: [source:STD-URBAN-DESIGN] for urban-design administration, [source:STD-CONTROL-PLAN] for the preparation and approval of detailed regulatory plans, [source:STD-LAND-USE] for territorial spatial land-and-sea-use classification, and [source:STD-ARCH-DEPTH-GAP] as a data-gap record for the required depth of architectural design documents. These materials are cited to constrain the depth of representation and define risk boundaries. They do not make this proposal a statutory plan, and a gap record cannot be treated as a verified provision.

Recent policy and site context comprises [source:BEIJING-AI-ORIGIN-2026], 2026 public information from the Beijing Municipal Commission of Development and Reform on Origin Community and the century-old Jing-Zhang corridor; [source:BEIJING-AI-DISTRICTS-2026], the “one core, multiple points” context for Beijing’s first AI innovation districts; [source:HAIDIAN-JZ-PHASE2-OPEN-2026], August 2026 reporting on the opening of Phase II; [source:HAIDIAN-AI-TRIAL-FIELD-2026], context on Haidian’s AI test field and Origin Community; [source:HAIDIAN-JZ-MIDTERM-2026], midterm outcomes and objectives for three priority groups and five-minute life support; and [source:HAIDIAN-15FYP-2026], district-level directions for AI, urban renewal, and public services in Haidian’s Fifteenth Five-Year Plan. The latter four come from this round of public-web desk research and support only a real-world baseline, policy intent, and questions for verification—not a site survey, as-built proof, parcel approval, or institutional commitment.

Governance and method context also includes [source:NATIONAL-DATA-INFRA-2025], the National Data Infrastructure Construction Guidelines; [source:AI-CONTENT-LABEL-2025], the Measures for Labeling AI-Generated and Synthetic Content; [source:AI-PLUS-2025], the State Council’s “AI+” opinion; [source:JZ-PARK-2023], public information on Phase I of the Jing-Zhang Railway Heritage Park; and [source:JZ-COCREATION-2021], material on public co-creation and professional continuity for the Jing-Zhang Railway Heritage Park. They support only direction, method, and governance boundaries; they create no project approval, funding, partnership, or control value.

The global comparison set includes [source:CASE-KENDALL], [source:CASE-ONE-NORTH], and [source:CASE-22AT]. The same set also includes [source:CASE-KINGS-CROSS], [source:CASE-STATION-F], and [source:CASE-MARS]. Existing directional context is [source:OSM-CONTEXT]. Its ODbL data is used only to identify roads, railways, and waterways and cannot serve as an official boundary or engineering basis.

Generation method: the spatial design is derived from the public provisional boundary, project enumerations, OSM context, and deterministic scripts. Areas and lengths are recalculated in EPSG:4548. Scenario passports, service zones, public-space modes, and phase gates are recorded as GeoJSON attributes. PNG, A3/A0, and offline HTML are explanatory layers, while geometry/*.geojson, metrics.json, the three matrices, and self_check.json are the evidence layer. COMMUNITY-DISPLAY-ONLY is recorded only as the current license label because its complete terms are not included. File-level status, unresolved items, and use limitations are stated in visual/assets/rights-clearance-ledger.json and report/copyright_statement.md.